

# Production-ready RAG helper classes

Package root used below: `com.example.rag`

Change the base package to match your project and add these files to `src/main/java` under the corresponding sub-packages.

## 1. `TieredApiVectorStoreCache.java`

```
package com.example.rag.cache;

import com.fasterxml.jackson.databind.JsonNode;
import com.fasterxml.jackson.databind.ObjectMapper;
import com.github.benmanes.caffeine.cache.Cache;
import com.github.benmanes.caffeine.cache.Caffeine;
import com.github.benmanes.caffeine.cache.LoadingCache;
import dev.langchain4j.data.document.Document;
import dev.langchain4j.model.embedding.EmbeddingModel;
import dev.langchain4j.model.embedding.LocalHuggingFaceEmbeddingModel;
import dev.langchain4j.store.embedding.EmbeddingStore;
import dev.langchain4j.store.embedding.EmbeddingStoreIngestor;
import dev.langchain4j.store.embedding.inmemory.InMemoryEmbeddingStore;
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;
import org.springframework.web.client.RestClientException;
import org.springframework.web.client.RestTemplate;

import java.time.Duration;
import java.util.List;
import java.util.stream.StreamSupport;

/**
 * Two-tier cache (short-lived burst level + long-lived level) that builds and
 * returns an {@link EmbeddingStore} for any REST endpoint.
 */
public class TieredApiVectorStoreCache {

    private static final Logger log =
        LoggerFactory.getLogger(TieredApiVectorStoreCache.class);

    private final RestTemplate rest;
    private final ObjectMapper mapper;
    private final EmbeddingModel embed;
```

```

private final Cache<String, EmbeddingStore<Document>> shortCache;
private final LoadingCache<String, EmbeddingStore<Document>> longCache;

public TieredApiVectorStoreCache(RestTemplate restTemplate,
                                ObjectMapper mapper,
                                long shortTtlSec,
                                long shortMaxEntries,
                                long longTtlMin,
                                long longMaxEntries) {
    this.rest = restTemplate;
    this.mapper = mapper;
    this.embed = LocalHuggingFaceEmbeddingModel.builder()
        .modelName("sentence-transformers/all-MiniLM-L6-v2")
        .build();

    this.shortCache = Caffeine.newBuilder()
        .expireAfterWrite(Duration.ofSeconds(shortTtlSec))
        .maximumSize(shortMaxEntries)
        .build();

    this.longCache = Caffeine.newBuilder()
        .expireAfterWrite(Duration.ofMinutes(longTtlMin))
        .maximumSize(longMaxEntries)
        .build(this::buildStore);
}

/**
 * Return a ready-to-search {@link EmbeddingStore}. Always non-null.
 */
public EmbeddingStore<Document> get(String url) {
    EmbeddingStore<Document> store = shortCache.getIfPresent(url);
    if (store == null) {
        store = longCache.get(url);           // may trigger loader
        shortCache.put(url, store);          // seed burst layer
    }
    return store;
}

/** Trigger an asynchronous rebuild for the given URL (used by scheduler).
 */
public void refresh(String url) {
    longCache.refresh(url);
    shortCache.invalidate(url);
    log.debug("Cache refresh triggered for {}", url);
}

```

```

/* ----- internal helpers ----- */

private EmbeddingStore<Document> buildStore(String url) {
    try {
        String json = rest.getForObject(url, String.class);
        List<Document> docs = toDocuments(url, json);

        EmbeddingStore<Document> store = new InMemoryEmbeddingStore<>();
        EmbeddingStoreIngestor.builder()
            .embeddingModel(embed)
            .embeddingStore(store)
            .build()
            .ingest(docs);

        log.info("Built vector store for {} ({} docs, {} B)", url,
docs.size(),
            json == null ? 0 : json.length());
        return store;
    } catch (RestClientException e) {
        log.warn("Failed to fetch {} - returning empty store", url, e);
        return new InMemoryEmbeddingStore<>();
    } catch (Exception ex) {
        throw new IllegalStateException("Cannot build vector store for " +
url, ex);
    }
}

private List<Document> toDocuments(String url, String json) throws
Exception {
    JsonNode root = mapper.readTree(json);
    JsonNode data = root.path("data").isMissingNode() ? root :
root.path("data");
    return StreamSupport.stream(data.spliterator(), false)
        .map(node -> Document.builder()
            .text(node.isContainerNode() ? node.toPrettyString() :
node.asText())
            .metadata("source", url)
            .build())
        .toList();
}
}

```

## 2. StaticEndpointReloader.java

```
package com.example.rag.cache;

import jakarta.annotation.PostConstruct;
import lombok.RequiredArgsConstructor;
import lombok.extern.slf4j.Slf4j;
import org.springframework.beans.factory.annotation.Value;
import org.springframework.scheduling.annotation.Scheduled;
import org.springframework.stereotype.Component;

import java.util.Map;

/**
 * Warms up & periodically refreshes static endpoints defined in YAML.
 */
@Slf4j
@Component
@RequiredArgsConstructor
public class StaticEndpointReloader {

    private final TieredApiVectorStoreCache cache;

    @Value("#{${rag.rest-api.static:}}}")
    private Map<String, String> staticEndpoints;

    @PostConstruct
    public void warmUp() {
        staticEndpoints.values().forEach(url -> {
            cache.get(url); // build & cache once
            log.info("Pre-loaded static endpoint {}", url);
        });
    }

    @Scheduled(cron = "${rag.rest-api.static-refresh-cron:0 0 * * * *}")
    public void refresh() {
        staticEndpoints.values().forEach(cache::refresh);
        log.debug("Static endpoints refreshed");
    }
}
```

Requires `spring-boot-starter` + `spring-boot-starter-aop` + `spring-boot-starter-scheduler` (enabled by `@EnableScheduling` in your application class).

## 2b. StaticDocumentReloader.java

```
package com.example.rag.cache;

import jakarta.annotation.PostConstruct;
import lombok.RequiredArgsConstructor;
import lombok.extern.slf4j.Slf4j;
import org.springframework.beans.factory.annotation.Value;
import org.springframework.scheduling.annotation.Scheduled;
import org.springframework.stereotype.Component;

import java.util.Map;

/**
 * Pre-loads static markdown / PDF paths at application startup and refreshes
 * them on a configurable cron schedule.
 */
@Slf4j
@Component
@RequiredArgsConstructor
public class StaticDocumentReloader {

    private final DocumentVectorStoreCache docCache;

    @Value("#{${rag.doc.static:{}}}")
    private Map<String, String> staticDocs; // name → path

    @PostConstruct
    public void warmUp() {
        staticDocs.values().forEach(p -> {
            docCache.get(p);
            log.info("Pre-loaded static doc {}", p);
        });
    }

    @Scheduled(cron = "${rag.doc.static-refresh-cron:0 0 * * * *}")
    public void refresh() {
        staticDocs.values().forEach(docCache::refresh);
        log.debug("Static documents refreshed");
    }
}
```

Like the endpoint reloader, this bean requires `@EnableScheduling` in your main application class.

### 3. MultiApiSimilarityRetriever.java

```
package com.example.rag.retriever;

import com.example.rag.cache.TieredApiVectorStoreCache;
import dev.langchain4j.data.document.Document;
import dev.langchain4j.model.embedding.EmbeddingModel;
import dev.langchain4j.rag.content.retriever.ContentRetriever;
import dev.langchain4j.rag.content.retriever.EmbeddingStoreContentRetriever;
import dev.langchain4j.rag.query.Query;
import dev.langchain4j.store.embedding.EmbeddingStore;
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;
import org.springframework.beans.factory.annotation.Value;
import org.springframework.stereotype.Component;

import java.util.*;

/**
 * Aggregates similarity search across configurable static endpoints and
 * per-request dynamic URLs supplied via Query metadata.
 */
@Component
public class MultiApiSimilarityRetriever implements ContentRetriever {

    private static final Logger log =
        LoggerFactory.getLogger(MultiApiSimilarityRetriever.class);

    private final Map<String, String> staticEndpoints; // name → URL
    private final TieredApiVectorStoreCache cache;
    private final EmbeddingModel embed;
    private final int topK;

    public MultiApiSimilarityRetriever(@Value("#{${rag.rest-api.static:{{{}}}")
        Map<String, String> staticEndpoints,
                                     TieredApiVectorStoreCache cache,
                                     EmbeddingModel embed,
                                     @Value("#{${rag.rest-api.top-k:6}") int
    topK) {
        this.staticEndpoints = staticEndpoints;
        this.cache = cache;
        this.embed = embed;
        this.topK = topK;
    }
}
```

```

@Override
public List<Document> retrieve(Query query) {
    Set<String> targetUrls = new HashSet<>();

    // 1 static endpoint selection
    endpointNamesFor(query).stream()
        .map(staticEndpoints::get)
        .filter(Objects::nonNull)
        .forEach(targetUrls::add);

    // 2 raw URLs via metadata("urls", "http://...")
    Object dyn = query.metadata("urls");
    if (dyn instanceof String s) {
        Arrays.stream(s.split(","))
            .map(String::trim)
            .filter(str -> str.startsWith("http"))
            .forEach(targetUrls::add);
    }

    if (targetUrls.isEmpty()) {
        log.debug("No endpoints selected - returning empty context");
        return List.of();
    }

    // 3 parallel similarity search per URL
    List<Document> combined = targetUrls.parallelStream()
        .flatMap(url -> {
            EmbeddingStore<Document> store = cache.get(url);
            return EmbeddingStoreContentRetriever.builder()
                .embeddingStore(store)
                .embeddingModel(embed)
                .maxResults(topK * 2)
                .build()
                .retrieve(query)
                .stream();
        })
        .distinct()           // de-dupe identical text
        .limit(topK)         // global top-k
        .toList();

    return combined;
}

private Set<String> endpointNamesFor(Query q) {
    Object meta = q.metadata("endpoints");
    if (meta instanceof String s) {
        return Arrays.stream(s.split(","))

```

```

        .map(String::trim)
        .filter(staticEndpoints::containsKey)
        .collect(Collectors.toSet());
    }
    return staticEndpoints.keySet();
}
}

```

#### 4. RagConfig.java

```

package com.example.rag.config;

import com.example.rag.cache.TieredApiVectorStoreCache;
import dev.langchain4j.model.embedding.EmbeddingModel;
import dev.langchain4j.model.embedding.LocalHuggingFaceEmbeddingModel;
import org.springframework.beans.factory.annotation.Value;
import org.springframework.context.annotation.Bean;
import org.springframework.context.annotation.Configuration;
import org.springframework.web.client.RestTemplate;
import com.fasterxml.jackson.databind.ObjectMapper;

@Configuration
public class RagConfig {

    /* Shared lightweight embedding model - loaded once per JVM */
    @Bean
    public EmbeddingModel localMiniLm() {
        return LocalHuggingFaceEmbeddingModel.builder()
            .modelName("sentence-transformers/all-MiniLM-L6-v2")
            .build();
    }

    @Bean
    public TieredApiVectorStoreCache vectorStoreCache(RestTemplate rest,
                                                    ObjectMapper mapper,
                                                    @Value("${rag.rest-api.short-ttl-seconds:90}") long shortTtl,
                                                    @Value("${rag.rest-api.short-max-entries:200}") long shortMax,
                                                    @Value("${rag.rest-api.long-ttl-minutes:60}") long longTtl,
                                                    @Value("${rag.rest-api.long-max-entries:100}") long longMax) {
        return new TieredApiVectorStoreCache(rest, mapper, shortTtl, shortMax,

```



```
longTtl, longMax);
    }
}
```

## 5. Example `application.yml` snippet

```
spring:
  main:
    allow-bean-definition-overriding: true

rag:
  rest-api:
    static:
      users:      https://reqres.in/api/users?page=2
      settings:   https://api.example.com/v1/settings
    top-k: 6
    parallelism: 4
    short-ttl-seconds: 90
    short-max-entries: 200
    long-ttl-minutes: 60
    long-max-entries: 100
    static-refresh-cron: "0 */30 * * * *" # every 30 minutes
```

### How to integrate

1. **Enable scheduling** in your Spring Boot application:

```
@SpringBootApplication
@EnableScheduling
public class RagApplication { public static void main(String[] args) {
    SpringApplication.run(RagApplication.class, args);
}}
```

2. **Add** the `MultiApiSimilarityRetriever` instance to your `LangChain4j Assistant`.
3. **Pass dynamic URLs** via `Query` metadata when needed:

```
Query q = Query.from(question)
    .withMetadata("urls", "https://api.example.com/user/42/
bookings");
```

You now have production-ready, pluggable RAG components that:

- pre-embed static endpoints at startup,
- refresh them on a cron,
- embed dynamic endpoints on demand,
- cache everything through a high-performance two-tier Caffeine strategy.

## 6. DocumentVectorStoreCache.java

```
package com.example.rag.cache;

import com.github.benmanes.caffeine.cache.*;
import com.fasterxml.jackson.databind.ObjectMapper;
import dev.langchain4j.data.document.Document;
import dev.langchain4j.data.document.loader.FileSystemDocumentLoader;
import dev.langchain4j.model.embedding.EmbeddingModel;
import dev.langchain4j.model.embedding.LocalHuggingFaceEmbeddingModel;
import dev.langchain4j.store.embedding.*;
import dev.langchain4j.store.embedding.inmemory.InMemoryEmbeddingStore;
import lombok.extern.slf4j.Slf4j;
import org.springframework.util.StringUtils;

import java.nio.file.Path;
import java.time.Duration;
import java.util.List;

/**
 * Two-tier Caffeine cache for *file-based* documents (Markdown, PDF, etc.).
 * Key is an absolute directory/file path.
 */
@Slf4j
public class DocumentVectorStoreCache {

    private final EmbeddingModel embed;
    private final Cache<String, EmbeddingStore<Document>> shortCache;
    private final LoadingCache<String, EmbeddingStore<Document>> longCache;

    public DocumentVectorStoreCache(long shortTtlSec, long shortMax,
                                    long longTtlMin, long longMax) {
        this.embed = LocalHuggingFaceEmbeddingModel.builder()
            .modelName("sentence-transformers/all-MiniLM-L6-v2")
            .build();

        this.shortCache = Caffeine.newBuilder()
            .expireAfterWrite(Duration.ofSeconds(shortTtlSec))
```

```

        .maximumSize(shortMax)
        .build();

    this.longCache = Caffeine.newBuilder()
        .expireAfterWrite(Duration.ofMinutes(longTtlMin))
        .maximumSize(longMax)
        .build(this::buildStore);
}

public EmbeddingStore<Document> get(String path) {
    EmbeddingStore<Document> store = shortCache.getIfPresent(path);
    if (store == null) {
        store = longCache.get(path);
        shortCache.put(path, store);
    }
    return store;
}

public void refresh(String path) {
    longCache.refresh(path);
    shortCache.invalidate(path);
}

/* ----- helpers ----- */
private EmbeddingStore<Document> buildStore(String pathStr) {
    try {
        Path path = Path.of(pathStr);
        List<Document> docs = FileSystemDocumentLoader.loadDocuments(path);

        EmbeddingStore<Document> store = new InMemoryEmbeddingStore<>();
        EmbeddingStoreIngester.builder()
            .embeddingModel(embed)
            .embeddingStore(store)
            .build()
            .ingest(docs);
        log.info("Built doc vector store for {} ({} docs)", path,
docs.size());
        return store;
    } catch (Exception e) {
        log.error("Failed to load documents from {}", pathStr, e);
        return new InMemoryEmbeddingStore<>();
    }
}
}

```

## 7. MultiDocSimilarityRetriever.java

```
package com.example.rag.retriever;

import com.example.rag.cache.DocumentVectorStoreCache;
import dev.langchain4j.data.document.Document;
import dev.langchain4j.model.embedding.EmbeddingModel;
import dev.langchain4j.rag.content.retriever.ContentRetriever;
import dev.langchain4j.rag.content.retriever.EmbeddingStoreContentRetriever;
import dev.langchain4j.rag.query.Query;
import dev.langchain4j.store.embedding.EmbeddingStore;
import org.springframework.beans.factory.annotation.Value;
import org.springframework.stereotype.Component;

import java.util.*;
import java.util.stream.Collectors;

/**
 * Similar to {@code MultiApiSimilarityRetriever} but for unstructured files.
 */
@Component
public class MultiDocSimilarityRetriever implements ContentRetriever {

    private final Map<String, String> staticPaths;           // name → dir/file
    path
    private final DocumentVectorStoreCache cache;
    private final EmbeddingModel embed;
    private final int topK;

    public MultiDocSimilarityRetriever(@Value("#{${rag.doc.static:}}}")
    Map<String,String> staticPaths,
                                     DocumentVectorStoreCache cache,
                                     EmbeddingModel embed,
                                     @Value("${rag.rest-api.top-k:6}") int
    topK) {
        this.staticPaths = staticPaths;
        this.cache = cache;
        this.embed = embed;
        this.topK = topK;
    }

    @Override
    public List<Document> retrieve(Query query) {
        Set<String> targetPaths = new HashSet<>();

        // static names from metadata("docs", "handbook,policies")
    }
```

```

Object meta = query.metadata("docs");
if (meta instanceof String s) {
    Arrays.stream(s.split(","))
        .map(String::trim)
        .filter(staticPaths::containsKey)
        .map(staticPaths::get)
        .forEach(targetPaths::add);
} else {
    targetPaths.addAll(staticPaths.values());
}

// raw dynamic path via metadata("path", "./uploads/123.md")
Object dyn = query.metadata("path");
if (dyn instanceof String s && !s.isBlank()) {
    targetPaths.add(s.trim());
}

return targetPaths.parallelStream()
    .flatMap(p -> {
        EmbeddingStore<Document> store = cache.get(p);
        return EmbeddingStoreContentRetriever.builder()
            .embeddingStore(store)
            .embeddingModel(embed)
            .maxResults(topK * 2)
            .build()
            .retrieve(query)
            .stream();
    })
    .distinct()
    .limit(topK)
    .toList();
}
}

```

## 8. CombinedContentRetriever.java

```

package com.example.rag.retriever;

import dev.langchain4j.data.document.Document;
import dev.langchain4j.rag.content.retriever.ContentRetriever;
import dev.langchain4j.rag.query.Query;

import java.util.*;
import java.util.stream.Collectors;

```

```

/**
 * Wraps multiple {@link ContentRetriever}s and merges their results.
 */
public class CombinedContentRetriever implements ContentRetriever {

    private final List<ContentRetriever> delegates;
    private final int topK;

    public CombinedContentRetriever(int topK, ContentRetriever... retrievers) {
        this.delegates = List.of(retrievers);
        this.topK = topK;
    }

    @Override
    public List<Document> retrieve(Query query) {
        return delegates.parallelStream()
            .flatMap(r -> r.retrieve(query).stream())
            .collect(Collectors.toMap(Document::text, d -> d, (a,b) -> a))
            .values()
            .stream()
            .limit(topK)
            .toList();
    }
}

```

## 9. Extended `RagConfig.java` additions

```

@Bean
public DocumentVectorStoreCache docVectorStoreCache(
    @Value("${rag.rest-api.short-ttl-seconds:90}") long sTtl,
    @Value("${rag.rest-api.short-max-entries:200}") long sMax,
    @Value("${rag.rest-api.long-ttl-minutes:60}") long lTtl,
    @Value("${rag.rest-api.long-max-entries:100}") long lMax) {
    return new DocumentVectorStoreCache(sTtl, sMax, lTtl, lMax);
}

@Bean
public CombinedContentRetriever combinedRetriever(MultiApiSimilarityRetriever
api,
                                                    MultiDocSimilarityRetriever
docs,
                                                    @Value("${rag.rest-api.top-k:
6}") int topK) {

```

```
    return new CombinedContentRetriever(topK, api, docs);  
}
```

## 10. YAML additions for documents

```
rag:  
  doc:  
    static:  
      handbook: classpath:docs/employee-handbook.md  
      policies:  ./legal/policies/  
      static-refresh-cron: "0 */60 * * * *" # hourly reload of markdown/pdfs
```

Now you have:

- `` — burst + long cache for any file path
- `` — similarity search over markdown/PDF paths
- `` — merges API + document hits and returns global top-k

Wire the new `combinedRetriever` into your `Assistant` instead of the individual ones, and you're ready to RAG across *both* REST data **and** unstructured files with one call.